

Analysis of IMDB Data

We will analyze a subset of IMDB's actors, genres, movie actors, and movie ratings data. This dataset comes to us from Kaggle (<https://www.kaggle.com/datasets/ashirwadsangwan/imdb-dataset> (<https://www.kaggle.com/datasets/ashirwadsangwan/imdb-dataset>)) although we have taken steps to pull this data into a public s3 bucket:

- s3://cis9760-lecture9-movieanalysis/name.basics.tsv ---> (actors)
- s3://cis9760-lecture9-movieanalysis/title.basics.tsv ---> (genres)
- s3://cis9760-lecture9-movieanalysis/title.principals.tsv ---> (movie actors)
- s3://cis9760-lecture9-movieanalysis/title.ratings.tsv ---> (movie ratings)

Content

name.basics.tsv.gz – Contains the following information for names:

nconst (string) - alphanumeric unique identifier of the name/person.
primaryName (string) – name by which the person is most often credited.
birthYear – in YYYY format.
deathYear – in YYYY format if applicable, else .
primaryProfession (array of strings) – the top-3 professions of the person.
knownForTitles (array of tconsts) – titles the person is known for.

title.basics.tsv.gz - Contains the following information for titles:

tconst (string) - alphanumeric unique identifier of the title.
titleType (string) – the type/format of the title (e.g. movie, short, tvseries, tvepisode, video, etc).
primaryTitle (string) – the more popular title / the title used by the filmmakers on promotional materials at the point of release.
originalTitle (string) - original title, in the original language.
isAdult (boolean) - 0: non-adult title; 1: adult title.
startYear (YYYY) – represents the release year of a title. In the case of TV Series, it is the series start year.
endYear (YYYY) – TV Series end year. for all other title types.
runtimeMinutes – primary runtime of the title, in minutes.
genres (string array) – includes up to three genres associated with the title.

title.principals.tsv – Contains the principal cast/crew for titles:

tconst (string) - alphanumeric unique identifier of the title.

ordering (integer) – a number to uniquely identify rows for a given titleId.
nconst (string) - alphanumeric unique identifier of the name/person.
category (string) - the category of job that person was in.
job (string) - the specific job title if applicable, else.
characters (string) - the name of the character played if applicable, else.

title.ratings.tsv.gz – Contains the IMDb rating and votes information for titles:

tconst (string) - alphanumeric unique identifier of the title.
averageRating – weighted average of all the individual user ratings.
numVotes - number of votes the title has received.

PART 1 - Installation and Initial Setup

Begin by installing the necessary libraries that you may need to conduct your analysis. At the very least, you must install pandas and matplotlib

In [1]:

```
%%info
```

```
Current session configs: {'conf': {'spark.pyspark.python': 'python3',  
'spark.pyspark.virtualenv.enabled': 'true',  
'spark.pyspark.virtualenv.type': 'native',  
'spark.pyspark.virtualenv.bin.path': '/usr/bin/virtualenv'},  
'kind': 'pyspark'}
```

No active sessions.

Let's install the necessary packages here

In [2]:

```
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
```

► Spark Job Progress

Starting Spark application

ID	YARN Application ID	Kind	State	
3	application_1669595400244_0004	pyspark	idle	Link

SparkSession available as 'spark'.

Collecting pandas==1.0.3

Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl (https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)

Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)

Using cached https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bbb257d0355c7f6128853c78955f57342a56d/python_dateutil-2.8.2-py2.py3-none-any.whl (https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bbb257d0355c7f6128853c78955f57342a56d/python_dateutil-2.8.2-py2.py3-none-any.whl)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)

Installing collected packages: python-dateutil, pandas

Successfully installed pandas-1.0.3 python-dateutil-2.8.2

Collecting matplotlib==3.2.1

Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3balf09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl (https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3balf09088b35776a799ba7dd95f7c2b195ec800933b276b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl)

Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1669611076034-0/lib/python3.7/site-packages (from matplotlib==3.2.1)

Collecting pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 (from matplotlib==3.2.1)

Using cached <https://files.pythonhosted.org/packages/6c/10/a7d0>

```
fa5baea8fe7b50f448ab742f26f52b80bfca85ac2be9d35cdd9a3246/pyparsing-3.0.9-py3-none-any.whl (https://files.pythonhosted.org/packages/6c/10/a7d0fa5baea8fe7b50f448ab742f26f52b80bfca85ac2be9d35cdd9a3246/pyparsing-3.0.9-py3-none-any.whl)
Collecting cyclr>=0.10 (from matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f51/cyclr-0.11.0-py3-none-any.whl (https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f51/cyclr-0.11.0-py3-none-any.whl)
Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)
Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/ab/8f/8dbe2d4efc4c0b08ec67d6efb7cc31fbfd688c80afad85f65980633b0d37/kiwisolver-1.4.4-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl (https://files.pythonhosted.org/packages/ab/8f/8dbe2d4efc4c0b08ec67d6efb7cc31fbfd688c80afad85f65980633b0d37/kiwisolver-1.4.4-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)
Collecting typing-extensions; python_version < "3.8" (from kiwisolver>=1.0.1->matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/0b/8e/f1a0a5a76cfef77e1eb6004cb49e5f8d72634da638420b9ea492ce8305e8/typing_extensions-4.4.0-py3-none-any.whl (https://files.pythonhosted.org/packages/0b/8e/f1a0a5a76cfef77e1eb6004cb49e5f8d72634da638420b9ea492ce8305e8/typing_extensions-4.4.0-py3-none-any.whl)
Installing collected packages: pyparsing, cyclr, typing-extensions, kiwisolver, matplotlib
Successfully installed cyclr-0.11.0 kiwisolver-1.4.4 matplotlib-3.2.1 pyparsing-3.0.9 typing-extensions-4.4.0
```

Now, import the installed packages from the previous block below.

In [3]:

```
import pandas as pd
import matplotlib as plt
```

Loading Data

Load all data from S3 into a Spark dataframe object

In [4]:

```
actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/name.basics.tsv',
genres = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.basics.tsv'
movie_actors = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.princ
movie_ratings = spark.read.csv('s3://cis9760-lecture9-movieanalysis/title.rati
```

► Spark Job Progress

Actors

Display the schema below:

In [5]:

```
actors.printSchema()
```

```
root
|-- nconst: string (nullable = true)
|-- primaryName: string (nullable = true)
|-- birthYear: string (nullable = true)
|-- deathYear: string (nullable = true)
|-- primaryProfession: string (nullable = true)
|-- knownForTitles: string (nullable = true)
```

Display the first 5 rows with the following columns:

- primaryName
- birthYear
- deathYear
- knownForTitles

In [6]:

```
actors.select("primaryName", "birthYear", "deathYear", "knownForTitles").show(5)
```

► Spark Job Progress

```
+-----+-----+-----+-----+
| primaryName | birthYear | deathYear | knownForTitles |
+-----+-----+-----+-----+
| Fred Astaire | 1899 | 1987 | tt0050419,tt00531... |
| Lauren Bacall | 1924 | 2014 | tt0071877,tt01170... |
| Brigitte Bardot | 1934 | \N | tt0054452,tt00491... |
| John Belushi | 1949 | 1982 | tt0077975,tt00725... |
| Ingmar Bergman | 1918 | 2007 | tt0069467,tt00509... |
+-----+-----+-----+-----+
only showing top 5 rows
```

Genres

Display the first 10 rows with the following columns:

- titleType
- primaryTitle
- genres

In [7]:

```
genres.select("titleType","primaryTitle","genres").show(10)
```

► Spark Job Progress

titleType	primaryTitle	genres
short	Carmencita	Documentary,Short
short	Le clown et ses c...	Animation,Short
short	Pauvre Pierrot	Animation,Comedy,...
short	Un bon bock	Animation,Short
short	Blacksmith Scene	Comedy,Short
short	Chinese Opium Den	Short
short	Corbett and Court...	Short,Sport
short	Edison Kinetoscop...	Documentary,Short
movie	Miss Jerry	Romance
short	Exiting the Factory	Documentary,Short

only showing top 10 rows

Display the unique categories below:

In [8]:

```
genres.select("titleType").distinct().show()
```

► Spark Job Progress

titleType
tvSeries
tvMiniSeries
movie
videoGame
tvSpecial
video
tvMovie
tvEpisode
tvShort
short

Display the schema below:

In [9]:

```
genres.printSchema()
```

```
root
|-- tconst: string (nullable = true)
|-- titleType: string (nullable = true)
|-- primaryTitle: string (nullable = true)
|-- originalTitle: string (nullable = true)
|-- isAdult: string (nullable = true)
|-- startYear: string (nullable = true)
|-- endYear: string (nullable = true)
|-- runtimeMinutes: string (nullable = true)
|-- genres: string (nullable = true)
```

Movie Actors

Display the schema below:

In [10]:

```
movie_actors.printSchema()
```

```
root
|-- tconst: string (nullable = true)
|-- ordering: string (nullable = true)
|-- nconst: string (nullable = true)
|-- category: string (nullable = true)
|-- job: string (nullable = true)
|-- characters: string (nullable = true)
```

Display the first 10 rows below

In [11]:

```
movie_actors.show(10)
```

► Spark Job Progress

	tconst	ordering	nconst	category
b	characters			jo
tt0000001	1	nm1588970	self	
\N ["Herself"]				
tt0000001	2	nm0005690	director	
\N				
tt0000001	3	nm0374658	cinematographer	director of phot
o...				
tt0000002	1	nm0721526	director	
\N				
tt0000002	2	nm1335271	composer	
\N				
tt0000003	1	nm0721526	director	
\N				
tt0000003	2	nm5442194	producer	produce
r				
tt0000003	3	nm1335271	composer	
\N				
tt0000003	4	nm5442200	editor	
\N				
tt0000004	1	nm0721526	director	
\N				
only showing top 10 rows				

Movie Ratings

Display the schema below:

In [12]:

```
movie_ratings.printSchema()
```

```
root
|-- tconst: string (nullable = true)
|-- averageRating: string (nullable = true)
|-- numVotes: string (nullable = true)
```

Display the first 10 rows in a descending order by the number of votes

In [13]:

```
from pyspark.sql.functions import col
movie_ratings.select("tconst", "averageRating", "numVotes").sort(col("numVotes
```

► Spark Job Progress

tconst	averageRating	numVotes
tt7430722	6.8	9999
tt4445154	8.1	9997
tt2229907	6.3	9996
tt0294097	8.0	9994
tt0264734	6.5	9993
tt8860450	6.3	9991
tt2032572	5.2	9991
tt0664505	8.4	999
tt7508752	7.9	999
tt1077089	7.3	999

only showing top 10 rows

Overview of Data

Display the number of rows and columns in each dataframe object.

In [14]:

```
print(f"Number of columns in Actors table: {len(actors.dtypes)}")
print(f"Number of rows in Actors table: {actors.count()}\n")

print(f"Number of columns in Genres table: {len(genres.dtypes)}")
print(f"Number of rows in Genress table: {genres.count()}\n")

print(f"Number of columns in Movie Actors table: {len(movie_actors.dtypes)}")
print(f"Number of rows in Movie Actors table: {movie_actors.count()}\n")

print(f"Number of columns in Movie Ratings table: {len(movie_ratings.dtypes)}")
print(f"Number of rows in Movie Ratings table: {movie_ratings.count()}\n")
```

► Spark Job Progress

```
Number of columns in Actors table: 6
Number of rows in Actors table: 9706922
```

```
Number of columns in Genres table: 9
Number of rows in Genress table: 6321302
```

```
Number of columns in Movie Actors table: 6
Number of rows in Movie Actors table: 36468817
```

```
Number of columns in Movie Ratings table: 3
Number of rows in Movie Ratings table: 993153
```

PART 2 - Analyzing Genres

Let's now answer this question: how many unique genres are represented in this dataset?

Essentially, we have the genres per movie as a list - this is useful to quickly see what each movie might be represented as but it is difficult to easily answer questions such as:

- How many movies are categorized as Comedy, for instance?
- What are the top 20 most popular genres available?

Association Table

We need to "break out" these genres from the tconst? One common approach to take is to build an association table mapping a single tconst multiple times to each distinct genre.

For instance, given the following:

tconst	titleType	genres
abcd123	XXX	a,b,c

We would like to derive something like:

tconst	titleType	genre
abcd123	XXX	a
abcd123	XXX	b
abcd123	XXX	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Implement the code necessary to derive the table described from the data set

In [15]:

```
genres.select("tconst", "titleType", "genres").show(5, truncate=False)
```

► Spark Job Progress

tconst	titleType	genres
tt0000001	short	Documentary, Short
tt0000002	short	Animation, Short
tt0000003	short	Animation, Comedy, Romance
tt0000004	short	Animation, Short
tt0000005	short	Comedy, Short

only showing top 5 rows

Display the first 10 rows of your association table below

In [16]:

```
import pyspark.sql.functions as F
genres_result = (genres
                  .withColumn('genres', F.explode(F.split('genres', ',')))
                  )

genres_result.select("tconst", "titleType", "genres").show(10)
```

► Spark Job Progress

```
+-----+-----+-----+
|  tconst|titleType|   genres|
+-----+-----+-----+
|tt0000001|   short|Documentary|
|tt0000001|   short|   Short|
|tt0000002|   short| Animation|
|tt0000002|   short|   Short|
|tt0000003|   short| Animation|
|tt0000003|   short|   Comedy|
|tt0000003|   short|   Romance|
|tt0000004|   short| Animation|
|tt0000004|   short|   Short|
|tt0000005|   short|   Comedy|
+-----+-----+-----+
only showing top 10 rows
```

Total Unique Genres

What is the total number of unique genres available in the movie category?

In [17]:

```
genres_result.filter(genres_result.titleType == "movie").select("genres").dist
```

► Spark Job Progress

29

What are the unique genres available?

In [18]:

```
genres_result.select("genres").distinct().show()
```

► Spark Job Progress

```
+-----+
| genres |
+-----+
| Mystery|
| Musical|
| Sport  |
| Action |
| Talk-Show|
| Romance|
| Thriller|
| \N     |
| Reality-TV|
| Family |
| Fantasy|
| History|
| Animation|
| Film-Noir|
| Short  |
| Sci-Fi |
| News   |
| Drama  |
| Documentary|
| Western|
+-----+
```

only showing top 20 rows

Oops! Something is off!

In [19]:

```
nll= '\\N'  
genres_result.select("genres").filter(col("genres") != nll).distinct().show()
```

► Spark Job Progress

```
+-----+  
| genres |  
+-----+  
| Mystery |  
| Musical |  
| Sport   |  
| Action  |  
| Talk-Show |  
| Romance |  
| Thriller |  
| Reality-TV |  
| Family  |  
| Fantasy |  
| History |  
| Animation |  
| Film-Noir |
```

In [20]:

```
nll= '\\N'  
genres_result.filter(genres_result.titleType == "movie").select("genres").filter
```

► Spark Job Progress

28

Top Genres by Movies

Now let's find the highest rated genres in this dataset by rolling up genres.

Average Rating / Genre

So now, let's unroll our distinct count a bit and display the per average rating value of per genre.

The expected output should be:

genre	averageRating
a	8.5

genre	averageRating
b	6.3
c	7.2

Or something to that effect.

First, let's join our two dataframes (movie ratings and genres) by tconst

In [21]:

```
genre = genres_result.join(movie_ratings, on=["tconst"], how = "inner")
genre.select("genres", "averageRating").filter(F.col("genres") != null).filter(r
```

► Spark Job Progress

```
+-----+-----+
|  genres | averageRating |
+-----+-----+
|   Drama |           4.2 |
|   Drama |           4.2 |
| Biography |           4.1 |
|   Drama |           4.1 |
|  History |           4.1 |
|   Drama |           5.7 |
|   Drama |           4.6 |
|  History |           4.6 |
| Biography |           6.3 |
|   Drama |           6.3 |
+-----+-----+
only showing top 10 rows
```

Now, let's aggregate along the averageRating column to get a resultant dataframe that displays average rating per genre.

In [22]:

```
nll = '\\N'
rating_and_genre.select("genres", "averageRating")\
    .withColumn("averageRating", F.col("averageRating").cast("float"))\
    .filter(F.col("genres") != nll)\
    .filter(F.col("titleType") == "movie")\
    .groupBy("genres").agg(F.avg("averageRating").alias("avg_rating"),)\
    .show(truncate=False)
```

► Spark Job Progress

genres	avg_rating
Mystery	5.940437537126316
Musical	6.203246053185319
Action	5.718734067904495
Sport	6.600145190943391
Talk-Show	5.800000190734863
Romance	6.125714179294426
Thriller	5.625967567519544
Reality-TV	6.379310377712907
Family	6.250560452699635
Fantasy	5.924820762891499
History	6.822718117193864
Animation	6.326203749467441
Film-Noir	6.636246780503378

Horizontal Bar Chart of Top Genres

With this data available, let us now build a barchart of all genres

HINT: don't forget about the matplotlib magic!

```
%matplotlib plt
```

In [23]:

```
nll = '\\N'
rating_and_genre.select("genres", "averageRating")\
    .withColumn("averageRating", F.col("averageRating").cast("float"))\
    .filter(F.col("genres") != nll)\
    .filter(F.col("titleType") == "movie")\
    .groupBy("genres").agg(F.avg("averageRating").alias("avg_rating"),)\
    .sort(F.desc("avg_rating"))\
    .show(truncate=False)
```

► Spark Job Progress

genres	avg_rating
Short	7.259999942779541
Documentary	7.245469805371099
News	7.200916040944689
Biography	6.983637643044585
Game-Show	6.974999904632568
History	6.822718117193864
Music	6.752020207214588
Film-Noir	6.636246780503378
Sport	6.600145190943391
War	6.483807036278403
Reality-TV	6.379310377712907
Animation	6.326203749467441
Drama	6.288080211097538
Family	6.250560452699635
Musical	6.203246053185319
Romance	6.125714179294426
Crime	6.026013333109149
Western	5.948970991005059
Comedy	5.941363107822231
Mystery	5.940437537126316

only showing top 20 rows

In [24]:

```
import matplotlib.pyplot as plt

nll = '\\N'
x = rating_and_genre.select("genres", "averageRating")\
    .withColumn("averageRating", F.col("averageRating").cast("float"))\
    .filter(F.col("genres") != nll)\
    .filter(F.col("titleType") == "movie")\
    .groupBy("genres").agg(F.avg("averageRating").alias("avg_rating"),)\
    .sort(F.asc("avg_rating")).toPandas().plot.barh(color='purple')
x.set_xlabel("Average Rating")
x.set_ylabel("Genre")
x.set_title("Top Genres in the Movie Category")


temp = rating_and_genre.select("genres", "averageRating")\
    .withColumn("averageRating", F.col("averageRating").cast("float"))\
    .filter(F.col("genres") != nll)\
    .filter(F.col("titleType") == "movie")\
    .groupBy("genres").agg(F.avg("averageRating").alias("avg_rating"),)\
    .sort(F.asc("avg_rating"))

labels = temp.select('genres').collect()
labels = [labels[i]["genres"] for i in range(len(labels))]

x.set_yticklabels(labels)
x.set_xlim([5, 7.5])

fig = plt.gcf()
fig.set_size_inches(18.5, 10.5)
%matplotlib plt
```

► Spark Job Progress



PART 3 - Analyzing Job Categories

Total Unique Job Categories

What is the total number of unique job categories?

In [25]:

```
movie_actors.select("tconst","category").distinct().show(5)
```

► Spark Job Progress

```
+-----+-----+
|  tconst | category |
+-----+-----+
|tt0000826|cinematographer|
|tt0001014|cinematographer|
|tt0001150|      actress  |
|tt0002234|      writer   |
|tt0002401|     director  |
+-----+-----+
only showing top 5 rows
```

In [26]:

```
movie_actors.select("category").distinct().count()
```

► Spark Job Progress

12

What are the unique job categories available?

In [27]:

```
movie_actors.select("category").distinct().show()
```

► Spark Job Progress

```
+-----+
|          category          |
+-----+
|          actress          |
|          producer         |
| production_designer       |
|          writer           |
|          actor            |
| cinematographer           |
|       archive_sound       |
|       archive_footage     |
|             self          |
|             editor         |
|          composer         |
|          director         |
+-----+
```

Top Job Categories

Now let's find the top job categories in this dataset by rolling up categories.

Counts of Titles / Job Category

The expected output should be:

category	count
a	15

category	count
b	2
c	45

Or something to that effect.

In [28]:

```
from pyspark.sql.functions import mean, stddev, col, abs, split, explode
job_categories = movie_actors.select('tconst', 'category').withColumn("category",
print(job_categories.groupBy('category').count().show())
```

► Spark Job Progress

```
+-----+-----+
|      category|  count|
+-----+-----+
|      actress|6325097|
|      producer|2197866|
|production_designer| 285924|
|       writer|4811596|
|       actor|8493701|
|cinematographer|1300404|
|  archive_sound|   2143|
|archive_footage| 209035|
|         self|6153089|
|       editor|1197669|
|      composer|1313187|
|      director|4179106|
+-----+-----+
```

None

Bar Chart of Top Job Categories

With this data available, let us now build a barchart of the top 5 categories.

HINT: don't forget about the matplotlib magic!

```
%matplotlib plt
```

In [29]:

```
job_categories_count = job_categories.groupBy('category').count()
job_categories_top_5 = job_categories_count.sort("count", ascending=False).lim
job_categories_top_5.show()
```

► Spark Job Progress


```
+-----+-----+
|category|  count|
+-----+-----+
|   actor|8493701|
| actress|6325097|
|   self |6153089|
|  writer|4811596|
|director|4179106|
+-----+-----+
```

In [30]:

```
job_categories_top_5_pd = job_categories_top_5.toPandas().set_index('category')

job_categories_top_5_pd.plot.bar(color='orange')
plt.title('Top Job Categories')
plt.ylabel("count")
plt.xlabel("categories")
plt.xticks(rotation = 45)
plt.ylim([3e6,9e6])
plt.tight_layout()
%matplotlib plt
```

► Spark Job Progress



PART 4 - Answer to the following questions:

1) Find all the "movies" featuring "Johnny Depp" and "Helena Bonham Carter".

First join actors, genres, and movie actors on each other

In [31]:

```
movie_actor_genre_join = movie_actors.join(genres, on=["tconst"],how = "inner")
move_actor_join = movie_actor_genre_join.join(actors, on=["nconst"],how = "inner")

actor_johnnyd = move_actor_join.select("primaryTitle","primaryName").filter(mov
actor_helena = move_actor_join.select("primaryTitle","primaryName").filter(mov

dfmix = actor_johnnyd.join(actor_helena, on=["primaryTitle"],how = "inner")
dfmix.select("primaryTitle").show(truncate=False)
```

► Spark Job Progress

```
+-----+
|primaryTitle|
+-----+
|Alice Through the Looking Glass|
|Dark Shadows|
|Charlie and the Chocolate Factory|
|Alice in Wonderland|
|Sweeney Todd: The Demon Barber of Fleet Street|
|Corpse Bride|
+-----+
```

2) Find all the "movies" featuring "Brad Pitt" after 2010.

In [32]:

```
nll = '\\N'
actor_braddp = move_actor_join.select("primaryTitle", "primaryName", "startYear")
actor_braddp.select("primaryTitle", "startYear").sort(F.desc("startYear")).show()
```

► Spark Job Progress

primaryTitle	startYear
Babylon	2021
Kajillionaire	2020
Irresistible	2020
Ad Astra	2019
The King	2019
Once Upon a Time ... in Hollywood	2019
Vice	2018
War Machine	2017
Voyage of Time: Life's Journey	2016
Allied	2016
By the Sea	2015
The Big Short	2015
Hitting the Apex	2015
Fury	2014
12 Years a Slave	2013
Kick-Ass 2	2013
World War Z	2013
Killing Them Softly	2012
Moneyball	2011
The Tree of Life	2011

3) What is the number of "movies" "acted" by "Zendaya" per year?

In [33]:

```
actor_zendaya = move_actor_join.select("primaryTitle", "primaryName", "startYear")
actor_zendaya.select('startYear')
print(actor_zendaya.groupBy('startYear').count().show())
```

► Spark Job Progress

```
+-----+-----+
|startYear|count|
+-----+-----+
|      2020|      1|
|      2018|      2|
|      2017|      1|
+-----+-----+
```

None

4) What are the "movies" by average rating greater than "9.7" and released in "2019"?

In [34]:

```
movie_rating_join = move_actor_join.join(movie_ratings, on=["tconst"], how = "inner")
ratings = movie_rating_join.select("primaryTitle", "startYear", "averageRating").\n\nratings.select("primaryTitle", "averageRating").sort(F.asc("averageRating")).show()
```

► Spark Job Progress

primaryTitle	averageRating
Kirket	10.0
The Butcher Baronet	10.0
A Medicine for the Mind	10.0
Bu Can Var Oldugu Sürece	10.0
Love in Kilnerry	10.0
A Grunt's Life	10.0
Our Scripted Life	10.0
L'Enfant Terrible	10.0
From Shock to Awe	9.8
Square One	9.8
Kamen Rider Zi-O: Over Quartzer	9.8
We Shall Not Die Now	9.8
Gini Helida Kathe	9.8
Time and motion	9.8
Randhawa	9.8
The Cardinal	9.9
Puritan: All of Life to The Glory of God	9.9
Superhombre	9.9

Extra Credit - Analysis of your choice

Try and analyze some interesting dimension to this data. You should specify the question in your Project2_Analysis.ipynb.

You must join at least two datasets.

In []:

